VIBE PROMPT

Deep Research Report

AI Prompt Perfector for Vibe Coders

Prepared February 2026 | For Non-Technical Founders

ABOUT THIS REPORT	*This report answers 8 key questions about building Vibe Prompt: a free web + mobile tool that transforms plain-language descriptions into perfect prompts for AI code editors (Claude Code, Cursor, Windsurf, Gemini CLI, Codex, OpenCode). All pricing and API limits reflect conditions as of February 2026 — flag: these change frequently.*

1. Competitor Table — Existing Prompt Tools

The market for prompt enhancement tools is growing fast. Here are the most relevant competitors Vibe Prompt will face, ranging from general enhancers to coding-specific tools:

Tool	Focus	Key Features	Pricing	Strengths	Weaknesses
PromptDC (promptdc.com)	Coding-first prompt rewriter	VS Code extension, Chrome extension, prompt history, JSON/YAML output modes, community prompts	Freemium (limited free)	Only tool built specifically for vibe coding / AI code editors; works inside editor	Small team, limited template library, less awareness
PromptPerfect (promptperfect.jina.ai)	General LLM prompt optimizer	Multi-model support (GPT-4, Midjourney), auto-optimization, API access	Free tier + ~$19.99/mo Pro	Polished UI, established user base, broad model support	Not coding-specific; feels generic for vibe coders
Cabina.AI (cabina.ai)	Multi-purpose prompt improver	Free without registration, code + writing + image modes, copy-paste workflow	Free (with account for full features)	No sign-up required to start, good for technical prompts	No coding-platform specificity; no template library
PromptEnhancer Pro (promptenhancer.pro)	Platform-specific enhancer	Separate tiles for Stable Diffusion, Midjourney, Code Projects; free to use	Currently free	Free, covers code generation as a category	Limited code depth; primarily image-focused
MaxAI (maxai.co)	Step-by-step prompt coach	Interactive feedback loop, iterative process, browser extension	Freemium	Teaches users to write better prompts over time	More of a coach than a transformer; slow workflow
16x Prompt (prompt.16x.engineer)	Context-rich coding prompt builder	Drag-and-drop file context, supports Claude/GPT-4, local data processing	Paid (one-time purchase)	Great for complex code tasks; privacy-focused (local processing)	Paid upfront; desktop app only; no mobile
Prompt Enhancer Pro (CodeDesign)	General marketing tool	Basic prompt enhancement as a side feature of a website builder	Bundled with paid plans	Part of a larger ecosystem	Not a standalone tool; not coding-focused at all

YOUR GAP	*No existing tool combines: (1) vibe coder focus, (2) free web + mobile access, (3) template library for specific tasks like auth/UI/debugging, and (4) beginner-friendly plain-language input. This is your moat.*

What Users Love vs. Hate About Existing Tools

Based on community reviews and user feedback gathered from Reddit, Product Hunt, and tool-specific communities:

What Users LOVE	What Users HATE
One-click transformation — describe it, get a perfect prompt instantly	Generic output that doesn't understand coding context
Seeing HOW their prompt improved (teaches prompting skills)	Tools that require accounts just to try the basic feature
Platform-specific outputs (Cursor vs. Windsurf vs. Claude Code differ)	No template library — users have to start from zero every time
Speed — fast transformation, no waiting	Paid walls for basic features when free alternatives exist
Saving and reusing successful prompts	Mobile-unfriendly interfaces (most tools are desktop-only)
Honest, specific technical language in the output	Vague or bloated prompts that confuse the AI code editor

2. Tech Stack Recommendation

As a non-technical founder, you need tools you can set up in days without writing code. Here are the best options, ranked by ease of use and how well they fit your specific use case:

Option A: Recommended — Webflow (Web) + Flutterflow or PWA (Mobile)

BEST FOR YOU	*Webflow gives you a beautiful, fast web tool. Pair it with a PWA (Progressive Web App — a website that works like a mobile app) to cover mobile without rebuilding everything. No app store needed initially.*

Layer	Tool	What It Does	Free Tier?	Monthly Cost (Paid)
Frontend (Web)	Webflow	Visual website builder with animations, forms, and custom code blocks. Your best choice for a polished web UI.	Yes (2 pages, webflow.io subdomain)	~$14–$23/mo
Frontend (Mobile)	PWA via Webflow	Webflow sites can be 'installed' on phones as Progressive Web Apps. No app store submission needed.	Yes (included with Webflow)	Free
Backend / Database	Supabase	Open-source Firebase alternative. Stores user prompts, saved templates, history. Has excellent free tier.	Yes (500MB DB, 50k MAU)	~$25/mo for Pro
AI Engine	Google Gemini 2.5 Flash API	Powers the prompt transformation. Best free tier available right now (250 RPD, 1M token context).	Yes — see Section 4	~$0.10/1M tokens paid
API Glue / Automation	Make (formerly Integromat)	Connects Webflow form submissions to your AI API and back, without coding.	Yes (1,000 ops/mo)	~$9/mo
Auth / User Accounts	Supabase Auth	Built into Supabase — email/OAuth login for free.	Yes	Included with Supabase
Hosting	Vercel or Webflow Hosting	Vercel for custom apps; Webflow hosting for Webflow sites.	Yes (both have generous free tiers)	~$20/mo

Option B: Quickest Launch — Bubble

Bubble is an all-in-one no-code platform where you can build both the web UI and the backend logic visually. It's slower and less polished than Webflow, but everything lives in one place. Good for testing your idea before investing in a more polished stack.

Pros	Cons
All-in-one: frontend + backend + database in one tool	Can't export your source code (locked in to Bubble forever)
Large plugin library — AI API connections exist	Pricing uses 'Workload Units' which are hard to predict
Large community, lots of tutorials	Not ideal for mobile (no native app publishing)
Free development tier	Slower performance than custom builds at scale

Option C: Fastest MVP — Bolt.new or v0.dev (Vibe Coding Your Own Tool)

META OPTION	*Use an AI code editor (like Bolt.new, Lovable, or v0.dev) to vibe code Vibe Prompt itself. Describe what you want, let the AI build a React app, deploy it to Vercel for free. This is the most technically ambitious option but ironically the fastest if you're comfortable prompting AI code tools.*

Tool	What It Is	Best For	Cost
Bolt.new	Full-stack AI app builder in browser — describe app, get deployable code	Fast MVP of web app with Supabase backend	Free tier, then ~$20/mo
Lovable (lovable.dev)	AI that builds React apps from prompts and deploys instantly	Polished UI without touching code	Free tier available
v0.dev (Vercel)	AI UI component generator — describe UI, get Tailwind/React code	Building beautiful UI components quickly	Free tier

Infrastructure Essentials — Free Tier Stack

Service	Purpose	Free Tier Details	Key Limit to Watch
Supabase	Database + Auth + Storage	500MB database, 50k monthly active users, 2GB file storage	Storage and bandwidth
Vercel	Hosting / Deployment	100GB bandwidth/mo, unlimited projects, custom domains	Function execution time
Cloudflare	CDN + DDoS protection	Unlimited bandwidth on CDN, free Workers (100k req/day)	Worker CPU time
Resend	Transactional email	3,000 emails/month free	Monthly send volume

Service	Purpose	Free Tier Details	Key Limit to Watch
GitHub	Code version control	Unlimited public repos, private repos for free users	Storage for large files

3. AI API Comparison — Powering the Prompt Engine

This is the heart of Vibe Prompt. You need an AI model to transform plain descriptions into structured prompts. Here's a plain-English comparison of every meaningful free option as of February 2026:

FLAG — FAST CHANGING	*AI API free tiers are extremely volatile. Google cut Gemini free limits by 50–92% in December 2025 with no warning. OpenAI removed its free API tier years ago. Always verify current limits at the provider's official documentation before committing to a stack.*

Provider	Best Free Model	Free Tier Limits (Feb 2026)	Credit Card Required?	Quality for Coding Prompts	Key Risks
Google Gemini (ai.google.dev)	Gemini 2.5 Flash	10 RPM, 250 RPD, 250K TPM, 1M token context window	No	Very Good — 1M context is unmatched for long codebases	EU/UK/CH users blocked on free tier. Limits cut 92% in Dec 2025 with no notice.
Groq (groq.com)	Llama 3.3 70B	~14,400 RPD, blazing fast (300+ tokens/sec)	No	Good — fast inference, solid for structured prompts	Rate limits less generous for heavy daily users. Model quality below GPT-4 class.
Mistral AI (mistral.ai)	Mistral 7B / Small	~1 billion tokens/month (experimental)	No	Good for structured output; strong at following formatting instructions	Free tier is experimental and may change. Lower reasoning quality than larger models.
OpenRouter (openrouter.ai)	30+ free models including DeepSeek R1, Llama 3.1 8B	50 RPD on free models (cut from 200 in April 2025)	No	Variable — choose model per task; DeepSeek R1 excellent for reasoning	Very low RPD (50/day) makes it unsuitable as primary API. Use as fallback only.
OpenAI	GPT-4o mini (via $5 trial credit)	$5 credit expires in 3 months; no ongoing free API	Yes (for API)	Excellent — gold standard for structured prompt generation	No sustainable free tier. Must pay. ~$0.15/1M input tokens for 4o mini.
Anthropic Claude	Claude Haiku 4.5	No free API tier (free chat interface only at claude.ai)	Yes (for API)	Excellent — best for nuanced instruction following	No free API access at all. Haiku starts at ~$0.80/1M tokens.

Provider	Best Free Model	Free Tier Limits (Feb 2026)	Credit Card Required?	Quality for Coding Prompts	Key Risks
Hugging Face (huggingface.co)	Various open models (Zephyr, Phi-3, etc.)	Inference API free with rate limits; Spaces free hosting	No	Variable — smaller open models struggle with complex prompt structure	Models vary wildly in quality. Free Inference API is slow and limited.

Our Recommendation: Multi-Provider Strategy

Don't rely on a single API. Build with a primary + fallback approach:

PRIMARY	*Google Gemini 2.5 Flash — Best balance of quality, context window, and free tier. No credit card. Use for the main transformation engine.*
FALLBACK #1	*Groq (Llama 3.3 70B) — Use when Gemini hits rate limits. Extremely fast. Switch to this when users need instant responses.*
FALLBACK #2	*Mistral Small — Use for volume-heavy scenarios. 1 billion tokens/month is a huge buffer for light prompts.*

When you have 100+ daily active users, move to a paid tier. At ~$0.10/1M tokens for Gemini Flash paid, even 10,000 prompt transformations/day (averaging 2,000 tokens each) costs roughly $2/day — very sustainable.

4. Template Library Blueprint

Templates are Vibe Prompt's secret weapon. Instead of users starting from scratch, they pick a category ('I want to build a login page') and the tool fills in the structural best practices automatically. Here are the 8 categories you should launch with, plus example prompts for each:

Category 1: UI & Frontend Components

WHO USES THIS	*Vibe coders building web or mobile interfaces who don't know how to describe responsive layouts, accessibility, or design systems to an AI.*

Example User Input:

"I want a login form that looks modern"

What Vibe Prompt Transforms It Into:

"Create a responsive login form component using [framework]. Requirements: (1) Email and password fields with proper input validation and error states. (2) Show/hide password toggle. (3) 'Remember me' checkbox. (4) 'Forgot password' link. (5) Submit button with loading state during API call. (6) Accessible labels and ARIA attributes. (7) Mobile-first responsive layout. (8) Match the existing design system (primary color: [color], border radius: [size]px). (9) Error messages should appear inline below the relevant field. Return a single, self-contained component with no external dependencies beyond [framework]."

Category 2: Database Design

WHO USES THIS	*Builders who know they need a database but can't express schema requirements, relationships, or indexing strategies in technical terms.*

Example User Input:

"I need a database for a task app where teams can assign tasks to people"

What Vibe Prompt Transforms It Into:

"Design a complete relational database schema for a team task management app using [database: PostgreSQL/Supabase]. Include: (1) Users table with id, email, name, avatar_url, created_at. (2) Teams table with id, name, owner_id (FK to users). (3) Team_members junction table with team_id, user_id, role (enum: owner/admin/member), joined_at. (4) Tasks table with id, title, description, status (enum: todo/in_progress/done), priority (enum: low/medium/high), assignee_id (FK to users, nullable), team_id (FK to teams), creator_id (FK to users), due_date, created_at, updated_at. (5) Include proper foreign key constraints and ON DELETE behavior. (6) Add indexes on frequently queried columns (assignee_id, team_id, status). (7) Provide the SQL CREATE TABLE statements and explain each relationship."

Category 3: Authentication & Authorization

WHO USES THIS	*Non-technical founders who understand they need 'login' but don't know how to specify role-based access, JWT flows, or OAuth.*

User Input: "Add Google login to my app"

Transformed: Full prompt specifying OAuth2 flow, callback URL handling, session management with Supabase/Auth.js, user profile creation on first login, error handling for failed auth, and security considerations including CSRF protection.

Category 4: API & Backend Logic

WHO USES THIS	*Builders who want their app to do something (send emails, process payments, connect to a third-party service) but can't describe it as an API endpoint.*

User Input: "When someone submits the contact form, send me an email"

Transformed: Full prompt specifying REST endpoint, input validation, rate limiting (5 submissions per IP per hour), email service integration (Resend/SendGrid), confirmation email to user, admin notification format, error responses, and logging requirements.

Category 5: Error Handling & Debugging

WHO USES THIS	*Vibe coders who got an error and just pasted it in — the template structures a proper debugging prompt with all context the AI needs.*

User Input: "TypeError: Cannot read property 'map' of undefined — fix this"

Transformed: A structured debugging prompt that includes the error message, asks for the relevant code context, specifies the likely causes to investigate in order (null check, async timing, API response shape), requests a fix with an explanation of why it works, and asks the AI to add a defensive check to prevent recurrence.

Category 6: Deployment & DevOps

Templates for: Vercel deployment config, environment variables setup, CI/CD with GitHub Actions, domain configuration, performance optimization (image optimization, caching headers), monitoring setup.

Category 7: Performance & Optimization

Templates for: Database query optimization, lazy loading, code splitting, bundle size analysis, caching strategy, Core Web Vitals improvement.

Category 8: Feature Addition to Existing Code

WHO USES THIS	*Vibe coders mid-project who need to add a specific feature without breaking what they've built. The template ensures they provide sufficient context.*

User Input: "Add dark mode to my app"

Transformed: A structured prompt asking the AI to (1) analyze the existing CSS/Tailwind setup, (2) implement a theme toggle with localStorage persistence, (3) use CSS variables or Tailwind's dark: prefix consistently, (4) handle system preference detection via prefers-color-scheme, (5) avoid flash of unstyled content on page load, and (6) provide the implementation without breaking existing light mode styles.

5. MVP Feature List — Must-Have vs. Nice-to-Have

Launch with the fewest features that deliver maximum value. Here's the prioritized breakdown:

Must-Have (Week 1 Launch)

Feature	Why It's Essential	Complexity to Build
Prompt transformation box	Core product — user types plain English, gets structured prompt back	Low (just an API call + display)
8 template categories (see Section 4)	Reduces friction — users pick a starting point	Medium (write good templates, wire to UI)
Copy-to-clipboard button	One-click copy is the key UX moment — make it frictionless	Low
Mobile-responsive web design	40%+ of users will try on phones first	Low (with Webflow or Tailwind)
Before/after display	Show original input vs. enhanced output side by side	Low
Free, no sign-up required	Eliminates the biggest barrier to first use	None (just don't require auth)
AI editor selector	Let user pick Cursor, Windsurf, Claude Code, etc. — output adapts	Medium (different system prompts per tool)

Nice-to-Have (Month 2–3)

Feature	Value	Priority
Saved prompt history (requires account)	Users love being able to re-use past prompts	High
Community prompt library	User-contributed templates — massive engagement driver	High
Prompt quality score / tips	Explains WHY the output is better — teaches users	Medium
Multi-step prompt chains	For complex tasks: 'First plan the architecture, then build it'	Medium
VS Code / Cursor extension	One-click prompt insertion without leaving the editor	High (but complex)
Dark mode	Table stakes for developer tools	Low
Export to .cursorrules / CLAUDE.md format	Power user feature for vibe coders who use rules files	Medium
API access for power users	B2B opportunity — let teams integrate Vibe Prompt in workflows	Low (for MVP)

DO NOT BUILD YET	*User collaboration, prompt marketplace (buying/selling), complex analytics dashboards, native iOS/Android apps, team workspaces. Build the core transformation tool first. Validate. Then expand.*

6. Development Roadmap — Step-by-Step Build Plan

This roadmap assumes you are a non-technical founder using AI code editors (you are building Vibe Prompt with vibe coding — very meta). Estimated timelines are realistic for someone using Claude Code, Cursor, or Lovable.

Phase 1: Proof of Concept (Days 1–3)

GOAL	*Build the simplest possible version that works. One text box, one button, one output. Share with 5 friends to validate the idea.*

- Day 1: Sign up for Google AI Studio (aistudio.google.com) and get your free Gemini API key. Copy it somewhere safe.
- Day 1: Open Bolt.new or Lovable and describe: 'Build a single-page web app with a text area for input, a dropdown for template categories, a button that calls the Gemini API to enhance the prompt, and displays the result in a styled output box with a copy button.'
- Day 2: Write your AI system prompt (see Section 7 for the exact framework). This is the secret sauce — the instructions you give to Gemini to transform prompts correctly.
- Day 2: Test with 20 real user inputs. Refine the system prompt until the output is reliably useful.
- Day 3: Deploy to Vercel (free). Share the URL with 5–10 vibe coders and collect feedback.

Phase 2: MVP Polish (Days 4–14)

GOAL	*Make it look professional enough to share publicly. Add the template library and mobile optimization.*

- Days 4–5: Move from Bolt.new to a more polished stack if needed, OR polish the existing UI. Add the 8 template categories as pre-filled starting points.
- Days 6–7: Build the 'AI editor selector' — when user picks Cursor vs. Claude Code vs. Windsurf, your system prompt adapts to each tool's quirks.
- Days 8–9: Mobile optimization. Test on iPhone and Android. Fix any layout issues.
- Days 10–11: Add optional Supabase backend for saving prompt history (optional auth flow — sign in to save, anonymous otherwise).
- Days 12–14: Add multi-provider fallback (Gemini primary, Groq backup). Handle rate limit errors gracefully with a user-friendly message.

Phase 3: Growth Foundation (Weeks 3–4)

GOAL	*Set up the infrastructure for organic growth: SEO, social sharing, and the community prompt library seed.*

- Week 3: Set up basic SEO. Each template category should have its own landing page (e.g., /templates/auth-prompts). Webflow makes this easy with dynamic CMS.
- Week 3: Create a 'Share this prompt' button that generates a shareable link — massive viral potential within the vibe coding community.
- Week 4: Soft launch to Reddit (r/vibecoding, r/cursor, r/ClaudeAI), Product Hunt (plan for a Tuesday–Thursday launch), and the Cursor/Windsurf Discord communities.
- Week 4: Set up simple analytics (Plausible Analytics — privacy-friendly, $9/mo or self-host free) to track which templates are most used.

Phase 4: Monetization Prep (Month 2)

Once you have 500+ regular users, you'll understand what to charge for. The most common path for this type of tool is a usage-based freemium:

- Free: 20 prompt transformations per day, 5 saved prompts, access to core templates
- Pro ($7–12/month): Unlimited transformations, unlimited saved prompts, community library access, priority AI model, VS Code extension
- Team ($25–49/month): Everything in Pro + team workspace, custom templates, API access

7. Prompt Engineering Patterns for Vibe Coding

This section is your competitive advantage. The system prompt you write to transform user inputs IS the product. Here are the proven patterns that work best for AI code editors specifically:

The Core Framework: SPEC Format

The best prompts for AI code editors follow a structure the vibe coding community calls SPEC:

SPEC Element	What It Is	Example
S — Situation/Context	What already exists in the codebase / what tech stack	"This is a React + Supabase app using Tailwind CSS and TypeScript"
P — Purpose/Goal	What the user wants to achieve (outcome, not implementation)	"Users need to be able to reset their password via email"
E — Explicit Constraints	Non-negotiables: frameworks, patterns, file structure, coding style	"Use Supabase Auth, follow the existing error handling pattern in /src/utils/errors.ts, no additional dependencies"
C — Checklist/Criteria	How to know when the task is done; edge cases to handle	"Handle: invalid email, expired token, already-used token. Return appropriate error messages for each."

Platform-Specific Patterns

Different AI code editors respond best to different prompt structures:

Tool	What It Does Well	Prompt Tips Specific to This Tool
Claude Code (CLI)	Long, complex multi-file tasks; great at following explicit rules files (CLAUDE.md)	Be verbose and explicit. Reference specific files with @filename. Ask it to 'plan first, then implement'. Claude Code loves structured markdown in prompts.
Cursor	Context-aware editing; best at understanding existing code via @-mentions	Always @-mention relevant files. Use Cursor's 'Plan mode' for complex tasks. Ask for 'explanation first' before code changes.
Windsurf (Cascade)	Deep context search; finds relevant files automatically; great for refactors	You don't need to @-mention files as much — Windsurf finds context itself. Focus on describing the outcome and letting Windsurf plan the implementation.
Gemini CLI	Large context window (1M tokens); great for	Use the context window advantage — include more files. Explicitly state which files should and shouldn't be modified.

Tool	What It Does Well	Prompt Tips Specific to This Tool
	whole-codebase tasks	
Lovable / Bolt.new	Full-stack generation from scratch; great for greenfield projects	Be extremely specific about UI — describe colors, layout, spacing. These tools need more visual guidance than code-editing tools.

The 5 Golden Rules Your System Prompt Should Enforce

- Rule 1: Always specify the tech stack explicitly — AI editors behave very differently when they know the framework
- Rule 2: Include a 'Definition of Done' — a checklist of edge cases and error states to handle
- Rule 3: Reference the existing code structure — tell the AI where things should go, not just what to build
- Rule 4: Separate 'planning' from 'implementation' for complex tasks — ask AI to confirm the plan before writing code
- Rule 5: Never use one giant prompt — break complex tasks into sequential steps

Your System Prompt Template (Use This to Build Vibe Prompt)

The following is the prompt you'll give to Gemini/Groq to power the transformation engine:

System Prompt Template for Vibe Prompt Engine
You are an expert prompt engineer specializing in AI-assisted software development. Your job is to transform a user's plain-language description into a detailed, structured prompt optimized for [SELECTED_AI_EDITOR]. The user's selected template category is: [TEMPLATE_CATEGORY] Rules for transformation: 1. Output ONLY the enhanced prompt. No explanation, no preamble, no 'Here is your enhanced prompt:' 2. Start with a clear role/context statement for the AI (e.g., 'You are a senior [framework] developer working on...') 3. Use the SPEC format: Situation → Purpose → Explicit Constraints → Checklist 4. Include specific error states and edge cases to handle 5. Specify output format (single component? multiple files? which files to modify?) 6. Add a quality check at the end: 'Before finishing, verify: [checklist]' 7. Adapt formatting for [SELECTED_AI_EDITOR] (Claude Code: markdown-heavy; Cursor: @-mention style; Windsurf: outcome-focused) 8. Keep the prompt under 800 words — detailed but not overwhelming User's input: [USER_INPUT]

8. Monetization — How Free Tools Sustain Themselves

Free tools have several proven paths to sustainability. Here's what works for developer-focused tools like Vibe Prompt:

Model	How It Works	Best For	Realistic Monthly Revenue (at 2,000 users)	Risk Level
Freemium	Free basic tier + paid plan for power features (unlimited usage, saved prompts, team features)	Developer tools where heavy users have clear ROI from upgrading	$1,000–5,000 (assuming 2–5% conversion at $10/mo)	Low — proven model
Usage Credits	Free 50 transformations/month, buy more ($5 for 500)	Tools where usage clearly correlates with value delivered	$500–2,000 (casual upsells)	Low
B2B / Team Plans	Sell to dev agencies, bootcamps, companies at $49–99/mo for team features	High LTV if you can reach team leads and technical managers	$2,000–10,000 (fewer but larger contracts)	Medium — requires outbound sales
Open Source + Donations	Open source the core, accept GitHub Sponsors / Ko-fi donations	Community trust, faster growth, contribution-based development	$100–1,000 (mostly from power users)	Medium — depends on community love
Affiliate / Referrals	Earn referral fees when users sign up for Claude, Cursor, Supabase via your links	Any tool that refers users to other paid services	$200–500 (passive)	Low — requires volume
Sponsored Templates	AI tool companies pay to have their product featured in relevant templates	Once you have 5,000+ users; Cursor, Windsurf, Supabase are natural sponsors	$500–3,000/mo per sponsor	Medium — audience required first

Recommended Monetization Path for Vibe Prompt

- Months 1–3: Stay completely free. Focus on growth. Launch on Product Hunt, Reddit, Twitter/X. Build community.
- Month 3: Add optional account creation with cloud-saved prompt history. Free accounts save up to 50 prompts. Start collecting emails.
- Month 4: Introduce Pro plan ($9/mo or $79/yr). Unlimited saves, team features, priority AI model, VS Code extension. Target: 2% of active users convert.
- Month 6: Add B2B outreach to coding bootcamps and dev agencies. Offer team plans.
- Month 6+: Explore sponsored integrations with Cursor, Windsurf, Supabase — natural partners whose users are your users.

9. Budget Breakdown

Full Free Tier Stack (Month 1)

TOTAL COST: $0/month	*You can launch and run Vibe Prompt at zero cost through the following free tiers, assuming under ~200 daily active users.*

Service	What For	Free Tier Limit	Cost
Google Gemini API	Prompt transformation engine	250 requests/day on Flash (10 RPM)	$0
Groq API	Fallback AI engine	~14,400 req/day (Llama 3.3 70B)	$0
Supabase	Database + auth + storage	500MB DB, 50k monthly active users	$0
Vercel	Hosting + deployment	100GB bandwidth/mo, unlimited projects	$0
Webflow (Free plan)	UI/website builder	2 pages, webflow.io subdomain	$0
Make (Free plan)	API automation glue	1,000 operations/month	$0
Plausible Analytics	Privacy-friendly analytics (self-hosted)	Free if self-hosted on Vercel	$0
GitHub	Version control	Unlimited free repos	$0
Cloudflare	CDN + domain DNS	Unlimited CDN bandwidth	$0
		TOTAL MONTHLY COST →	$0

Paid Stack (When You Have 500+ Daily Active Users)

TOTAL COST: ~$80–150/month	*At scale, you'll need paid API access and more generous hosting. This budget covers up to ~5,000 daily users, at which point revenue should more than cover costs.*

Service	Why You Need Paid	Estimated Cost
Gemini API (Paid Tier 1)	250 req/day free isn't enough at scale. Paid: ~$0.10/1M tokens for Flash	~$10–30/mo (at 100k daily tokens)
Supabase Pro	More storage, better performance, no pause risk	$25/mo
Vercel Pro	Better performance, more bandwidth, team features	$20/mo

Service	Why You Need Paid	Estimated Cost
Webflow CMS/Business	Needed for dynamic template pages + custom domain	$23–36/mo
Make (Core plan)	More automation operations for user workflows	$9/mo
Domain name	yourdomain.com (annual)	~$12/year (~$1/mo)
	TOTAL MONTHLY COST →	~$88–121/mo

Revenue vs. Cost Breakeven

Scenario	Monthly Users	Pro Conversions (2%)	Revenue at $9/mo	Costs	Profit
Early MVP	500	10 users	$90	$0 (free tier)	$90
Growing	2,000	40 users	$360	~$90/mo	$270
Sustainable	5,000	100 users	$900	~$120/mo	$780
Success	20,000	400 users	$3,600	~$300/mo	$3,300

Quick Reference — Key Links & Resources

Tools to Sign Up For Today (All Free)

Service	URL	What You'll Do There
Google AI Studio	aistudio.google.com	Get your Gemini API key
Groq Console	console.groq.com	Get your backup Groq API key
Supabase	supabase.com	Set up your free database + auth
Vercel	vercel.com	Deploy your app for free
Bolt.new	bolt.new	Vibe code your MVP
Lovable	lovable.dev	Alternative to Bolt — AI builds the UI
OpenRouter	openrouter.ai	Access 30+ free AI models (fallback)
Webflow	webflow.com	Build a polished web frontend

Communities to Join for Research & Launch

- r/vibecoding — The primary Reddit community for vibe coders; your target users
- r/cursor — Cursor IDE users; huge overlap with your audience
- r/ClaudeAI — Claude users and Claude Code enthusiasts
- Cursor Discord (discord.gg/cursor) — Highly active developer community
- Windsurf Community Discord — Another major vibe coder hub
- Twitter/X: Follow @karpathy (coined 'vibe coding'), @anthropic, @cursor_ai for trends
- Product Hunt — Launch here on a Tuesday/Wednesday for max visibility

Competing Tools to Monitor

- PromptDC (promptdc.com) — Your closest direct competitor; monitor their product updates
- PromptPerfect (promptperfect.jina.ai) — General prompt optimization market leader
- Cursor's built-in prompt improvement features — Cursor may build this natively
- Anthropic's prompt engineering docs (docs.anthropic.com) — Free best practices

FINAL NOTE	*The vibe coding market is moving extremely fast. Andrej Karpathy coined the term in early 2025 and it's already a major development paradigm by early 2026. You have a real window to establish Vibe Prompt as the go-to prompt perfection tool for this community before the AI code editors build it natively. Move fast, stay free, and listen obsessively to your early users.*

— End of Report —